

OS Project 2

Scheduling in Linux

Advisor: Tei-Wei Kuo

Speaker: Chun-Feng Wu

Yu-Chen Lin

Outline

- ▶ Review: Life of Process
- ▶ Scheduling in Linux
- ▶ Implementation
 - ▶ Part I: Invoke FIFO Scheduler
 - ▶ Part II: Weighted Round Robin Scheduler
- ▶ Submission Rules
- ▶ References

Outline

- ▶ Review: Life of Process
- ▶ Scheduling in Linux
- ▶ Implementation
 - ▶ Part I: Invoke FIFO Scheduler
 - ▶ Part II: Weighted Round Robin Scheduler
- ▶ Submission Rules
- ▶ References

Process Life Cycle

- ▶ A process is **not** always ready to run.
- ▶ The scheduler must know the status of every process in the system when switching between tasks.
- ▶ A process may have one of the following **states**:
 - ▶ **Running** — The process is **executing at the moment**.
 - ▶ **Waiting** — The process is able to run but is not allowed to **because the CPU is allocated to another process**. The scheduler can select the process at the next task switch.
 - ▶ **Sleeping** — The process is sleeping and cannot run because it is **waiting for an external event**. The scheduler cannot select the process at the next task switch.
- ▶ The system saves all processes in a process table.

The Need of the Scheduler

- ▶ A unique description of each process is held in memory and is linked with other processes by means of several structures.
- ▶ This is the situation facing the scheduler, whose task is **to share CPU time between the programs to create the illusion of concurrent execution.**
- ▶ This task is split into two different parts —
 - ▶ One relating to the **scheduling policy** and
 - ▶ The other to context switching

Outline

- ▶ Review: Life of Process
- ▶ Scheduling in Linux
- ▶ Implementation
 - ▶ Part I: Invoke FIFO Scheduler
 - ▶ Part II: Weighted Round Robin Scheduler
- ▶ Submission Rules
- ▶ References

Scheduling in Linux (1/2)

- ▶ The schedule function is the starting point to an understanding of scheduling operations.
- ▶ It is defined in “`kernel/sched.c`” and is one of the most frequently invoked functions in the kernel code.
- ▶ Not only priority scheduling but also two other soft real-time policies required by the POSIX standard are implemented.
 - ▶ E.g., completely fair scheduling, real-time scheduling and scheduling of the idle task, etc.

Scheduling in Linux (2/2)

- ▶ The scheduler uses a series of data structures to sort and manage the processes in the system.
- ▶ Scheduling can be activated in two ways:
 - ▶ **Main scheduler**: Either directly if a task goes to sleep or wants to yield the CPU for other reasons,
 - ▶ **Periodic scheduler**: Or by a periodic mechanism that is run with constant frequency to check from time to time if switching tasks is necessary
- ▶ **Generic scheduler = Main + Periodic schedulers**

Generic Scheduler

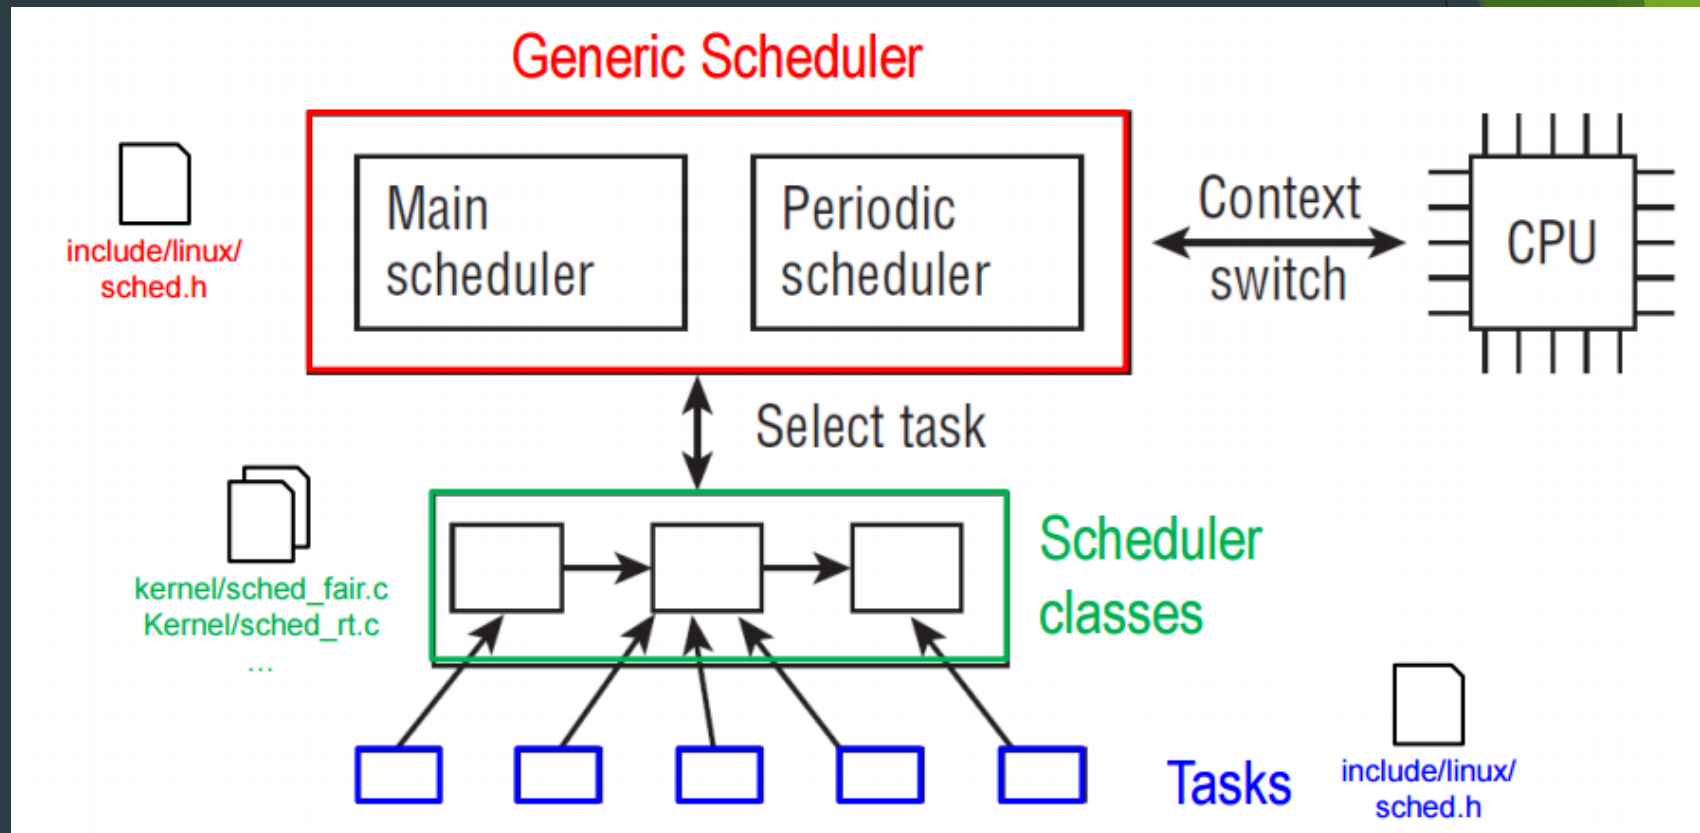
Scheduler Classes

Task

Task

Task

Overview of the Scheduling Subsystem in Linux



Task Representation

- In Linux, all concerned with processes and programs are built around a data structure: **task_struct**.

<sched.h>

```
struct task_struct {  
    volatile long state;      /* -1 unrunnable, 0 runnable, >0 stopped */  
    void *stack;  
    atomic_t usage;  
    unsigned long flags;      /* per process flags, defined below */  
    unsigned long ptrace;  
    int lock_depth;           /* BKL lock depth */  
  
    int prio, static_prio, normal_prio;  
    struct list_head run_list;  
    const struct sched_class *sched_class;  
    struct sched_entity se;  
};
```

... see more in "include/linux/sched.h"

Policy Designation

```
<sched.h>
struct task_struct {
    ...
    int prio, static_prio, normal_prio;
    unsigned int rt_priority;

    struct list_head run_list;
    const struct sched_class *sched_class;
    struct sched_entity se;

    unsigned int policy;
    cpumask_t cpus_allowed;
    unsigned int time_slice;

    ...
}
```

Scheduler Classes (1 / 3)

- ▶ Scheduler classes provide the connection between the **generic scheduler** and **individual scheduling methods**.
 - ▶ They are represented by several **function pointers** collected in a special data structure.
 - ▶ Each operation that can be requested by the global scheduler is represented by one pointer.
- ▶ This allows for creation of the generic scheduler **without** any knowledge about the internal working of different scheduler classes.

Scheduler Classes (2/3)

- An instance of **struct sched_class** must be provided for each scheduling class.

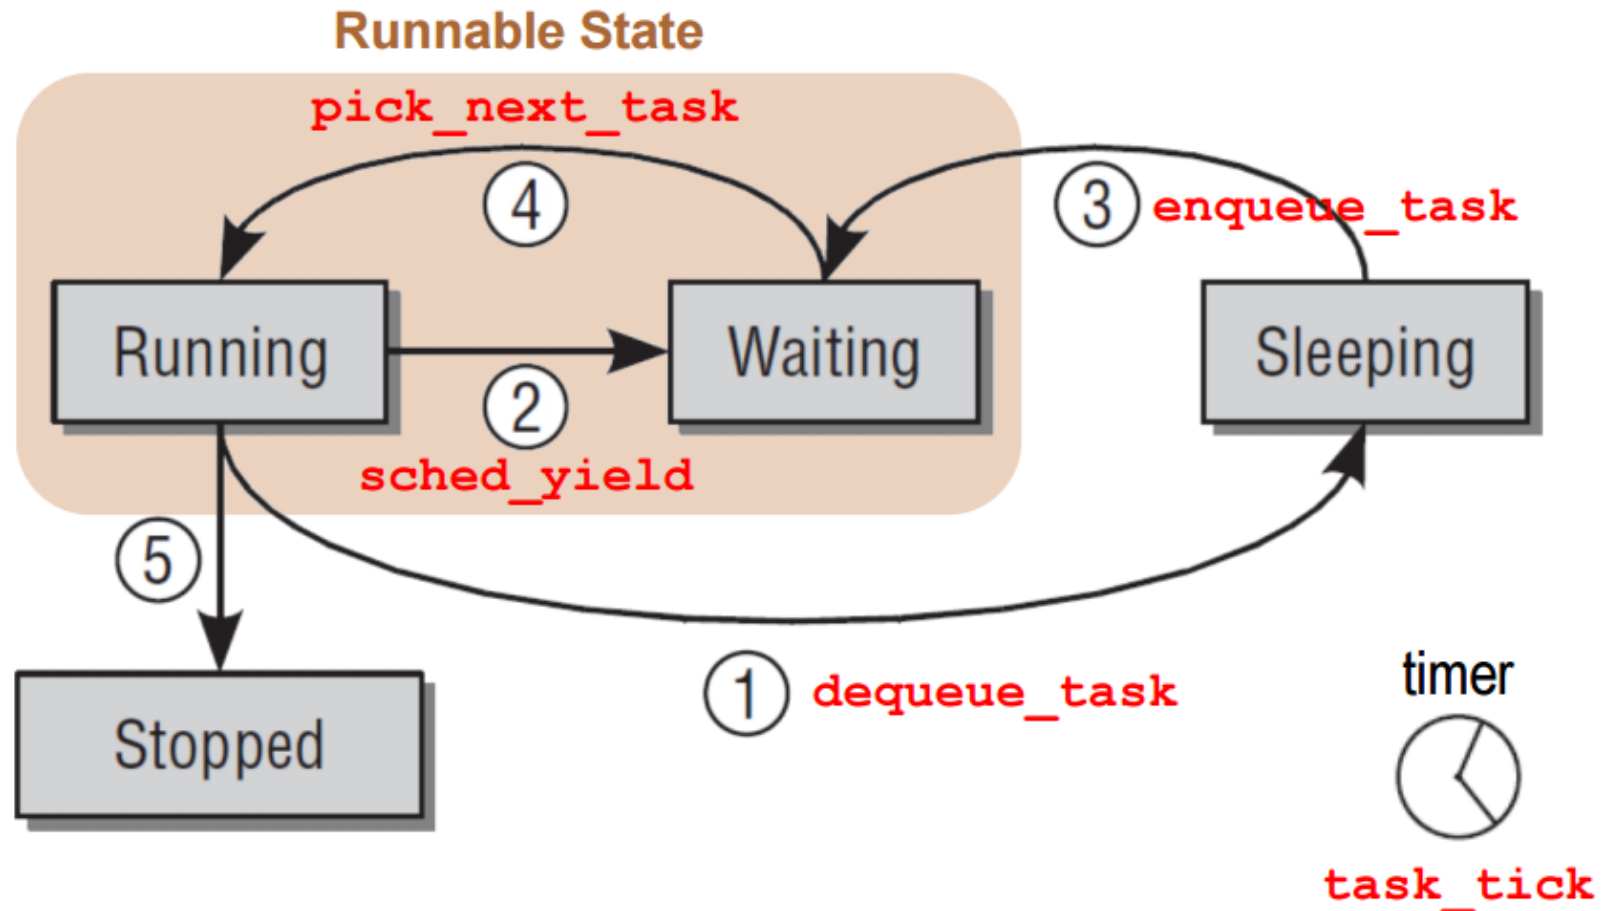
```
<sched.h>
```

```
struct sched_class {  
    const struct sched_class *next;  
  
    void (*enqueue_task) (struct rq *rq, struct task_struct *p, int wakeup);  
    void (*dequeue_task) (struct rq *rq, struct task_struct *p, int sleep);  
    void (*yield_task) (struct rq *rq);  
  
    void (*check_preempt_curr) (struct rq *rq, struct task_struct *p);  
  
    struct task_struct * (*pick_next_task) (struct rq *rq);  
    void (*put_prev_task) (struct rq *rq, struct task_struct *p);  
    void (*set_curr_task) (struct rq *rq);  
    void (*task_tick) (struct rq *rq, struct task_struct *p);  
    void (*task_new) (struct rq *rq, struct task_struct *p);  
};
```

Scheduler Classes (3/3)

- ▶ **enqueue_task**: adds a new process to the run queue. This happens when a process changes from a sleeping into a runnable state.
- ▶ **dequeue_task**: provides the inverse operation: It takes a process off a run queue. Naturally, this happens when a process switches from a runnable into an un-runnable state, or when the kernel decides to take it off the run queue for other reasons.
- ▶ **yield_task** : when a process wants to **relinquish control of the processor voluntarily**, it can use the sched_yield system call. This triggers yield_task to be called in the kernel.
- ▶ **pick_next_task**: selects the next task that is supposed to run
- ▶ **task_tick**: is called by the **periodic scheduler** each time it is activated.

Relationships between Generics Functions and Process States



Outline

- ▶ Review: Life of Process
- ▶ Scheduling in Linux
- ▶ Implementation
 - ▶ Part I: Invoke FIFO Scheduler
 - ▶ Part II: Weighted Round Robin Scheduler
- ▶ Submission Rules
- ▶ References

Linux Scheduling Policies

- ▶ Linux Scheduling Policies
 - ▶ Normal Scheduling policies (Non-real-time)
 - ▶ SCHED_OTHER, SCHED_BATCH, SCHED_IDLE.
 - ▶ Real-Time policies
 - ▶ SCHED_FIFO, SCHED_RR.
- ▶ The default scheduling policy is **non-real-time**.
- ▶ In this part, using Linux **real-time scheduling policy (FIFO)** to schedule threads in a process.

Part I: Invoke FIFO Scheduler

- ▶ Write a C program (sched_test.c) to create two threads.
- ▶ Each thread will print who is running and busy for 0.5 second.
- ▶ Run the program by default time-sharing schedule policy and show the result.

Ex. `$./sched_test`

- ▶ Run the program by real-time scheduling policy (FIFO) and show the result.

Ex. `$./sched_test SCHED_FIFO`

```
1  int main() {
2      set CPU affinity//all threads run on the same core
3      invoke FIFO_SCHED
4
5      for(i=0;i<2;i++)
6          thread_create(i)
7          print "Thread i was created"
8      for(i=0;i<2;i++)
9          thread_join(i)
10 }
11 thread_func() {
12     for(i=0;i<3;i++)
13         print "Thread # is running"
14         busy 0.5 second
15 }
```

Result

```

#( 04/17/17@10:34PM )( cfwu@ubuntu ):~
  sudo ./sched_test
Thread 1 was created
Thread 2 was created
Thread 2 is running
Thread 1 is running
Thread 2 is running
Thread 1 is running
Thread 1 is running
Thread 2 is running
#( 04/17/17@10:34PM )( cfwu@ubuntu ):~
  sudo ./sched_test SCHED_FIFO
Thread 1 was created
Thread 2 was created
Thread 1 is running
Thread 1 is running
Thread 1 is running
Thread 2 is running
Thread 2 is running
Thread 2 is running

```

Hint for Part I

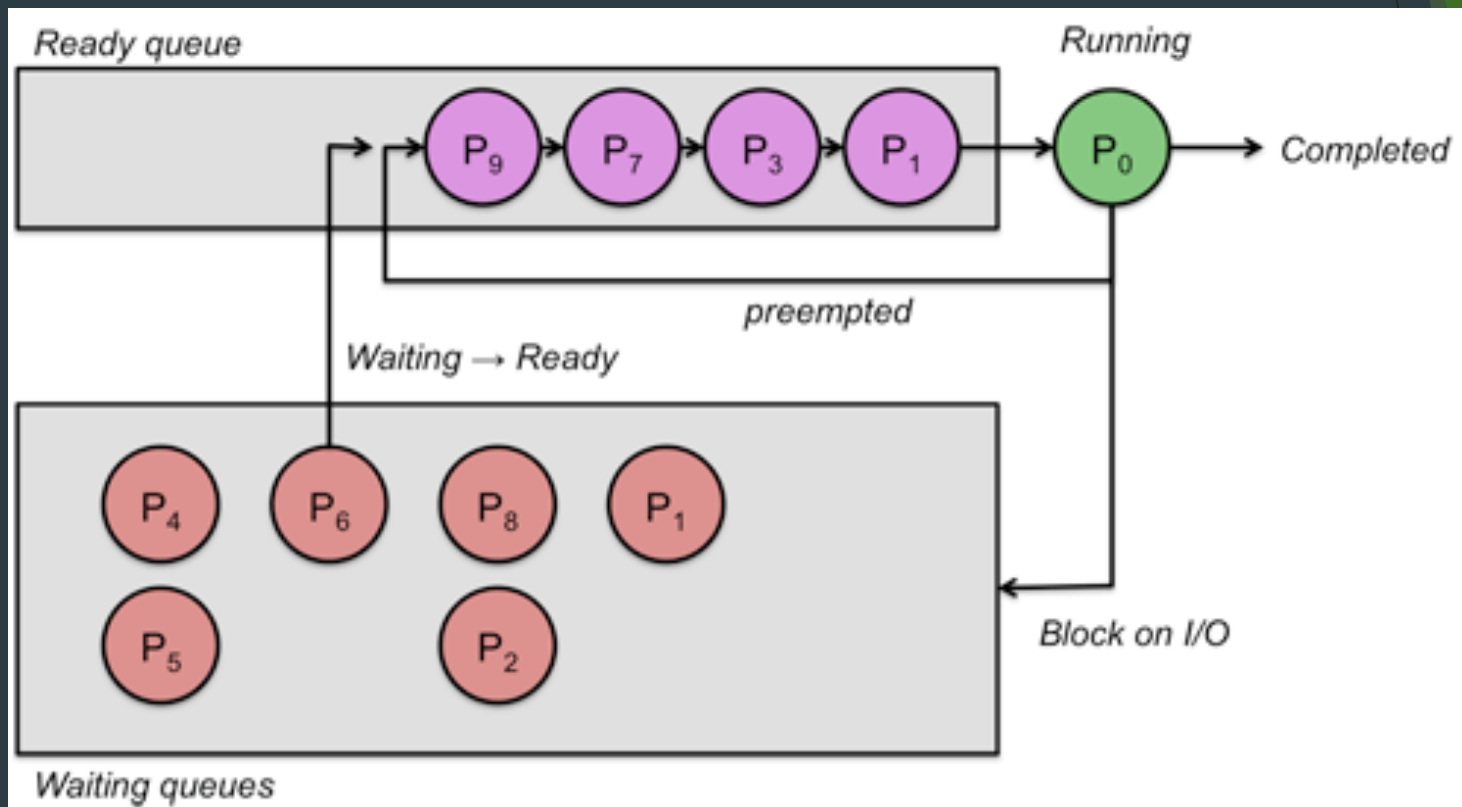
- ▶ Set CPU affinity
- ▶ `sched_setscheduler();`
- ▶ The policy corresponding value define in `/include/linux/sched.h`
- ▶ Set the priority of real-time process (`sched_param *param`)
- ▶ The permission to run real-time process

Outline

- ▶ Review: Life of Process
- ▶ Scheduling in Linux
- ▶ Implementation
 - ▶ Part I: Invoke FIFO Scheduler
 - ▶ Part II: Weighted Round Robin Scheduler
- ▶ Submission Rules
- ▶ References

Part II: Weighted Round Robin Scheduling (1/2)

- Processes are dispatched in a FIFO sequence but each process is allowed to run for only a limited amount of time, a.k.a., **time-slice** or **quantum**.



Part II: Weighted Round Robin Scheduling (2/2)

- ▶ Implement `kernel/sched_weighted_rr.c`
 - ▶ `enqueue_task_weighted_rr()`
 - ▶ `dequeue_task_weighted_rr()`
 - ▶ `yield_task_weighted_rr()`
 - ▶ `pick_next_task_weighted_rr()`
 - ▶ `task_tick_weighted_rr()`
- ▶ Tasks with higher weights can finish their jobs earlier by having larger time slices.

How to add a custom scheduler
into Linux?

Generic Scheduler Side (1/3)

In “include/linux/sched.h”,

- Add **#define SCHED_WEIGHTED_RR 6** - to define your weighted rr policy

```
32  /*
33   * Scheduling policies
34   */
35  #define SCHED_NORMAL          0
36  #define SCHED_FIFO            1
37  #define SCHED_RR              2
38  #define SCHED_BATCH           3
39  /* SCHED_ISO: reserved but not implemented yet */
40  #define SCHED_IDLE            5
41  //+ OS Proj2: weighted_rr
42  #define SCHED_WEIGHTED_RR     6
```

Generic Scheduler Side (2/3)

In “kernel/sched.c”

- Modify `__setscheduler()`, and `__sched_setscheduler()` functions - to let the generic scheduler can recognize your weighted rr scheduler

```
6515      //+ OS Proj2: weighted_rr
6516      case SCHED_WEIGHTED_RR:
6517          p->sched_class = &weighted_rr_sched_class;
6518          break;
6519      }
```

Generic Scheduler Side (3/3)

In “struct rq” of “kernel/sched.c”

- Add **struct weighted_rr_rq weighted_rr** - to specify the run queue for your weighted rr

```
506 struct rq {  
507     ...  
508     struct cfs_rq cfs;  
509     //+ OS Proj2: weighted_rr  
510     struct weighted_rr_rq weighted_rr;  
511     struct rt_rq rt;  
512     ...  
}
```

Note that struct rq - the generic per-CPU run queue structure. However, this is **NOT** the queue structure you will work with. Rather, this structure contains a more specific run queue type for different scheduler classes.

Scheduler Classes Side (1/3)

As well in “kernel/sched.c”

- ▶ Define `weighted_rr_rq` structure, which should contain
 - ▶ `struct list_head queue` - to denote the actual run queue for your weighted rr scheduler
 - ▶ `unsigned long nr_running` - to denote the number of processes which are now in the run queue

```
424 //+ OS Proj2: weighted_rr
425 struct weighted_rr_rq {
426     struct list_head queue;
427     unsigned long nr_running;
```

Scheduler Classes Side (2/3)

In “kernel/sched.c”,

- Declare `int weighted_rr_time_slice` - to define the time slice for your weighted rr scheduling policy

```
1934 //+ OS Proj2: weighted_rr
1935 int weighted_rr_time_slice
```

```
7227 //+ OS Proj2: weighted_rr
7228 SYSCALL_DEFINE1(sched_weighted_rr_setquantum, unsigned int, quantum)
7229 {
7230     weighted_rr_time_slice = quantum;
7231     return;
7232 }
```

Scheduler Classes Side (3/3)

In “kernel/sched_weighted_rr.c”

- ▶ Accomplish the implementation of weighted rr scheduler
 - ▶ Recall that an instance of **struct sched_class** must be provided for each scheduling class.

```
243  const struct sched_class weighted_rr_sched_class = {
244      .next                = &idle_sched_class,
245      .enqueue_task        = enqueue_task_weighted_rr,
246      .dequeue_task        = dequeue_task_weighted_rr,
247      .yield_task          = yield_task_weighted_rr,
248
249      .check_preempt_curr  = check_preempt_curr_weighted_rr,
250
251      .pick_next_task       = pick_next_task_weighted_rr,
252      .put_prev_task       = put_prev_task_weighted_rr,
253      ...
```

Task Side

In “struct task_struct” of “include/linux/sched.h”, add

- ▶ Declare **unsigned int weighted_rr_task_time_slice** - to denote the current time slice for this task
- ▶ Declare **struct list_head weighted_rr_list_item** - to denote the list item which will be inserted into the run queue of weighted_rr

```
1219 struct task_struct {
1220     ...
1221     //+ OS Proj2: weighted_rr
1222     unsigned int task_time_slice;
1223     unsigned int weighted_time_slice;
1224     ...
1225     //+ OS Proj2: weighted_rr
1226     struct list_head weighted_rr_list_item;
```

“Lazy Package”

- ▶ The lazy package includes

- ▶ <http://newslab.csie.ntu.edu.tw/course/OS2017/files/project/PJ2-linux-2.6.32.60.tar.gz>

- ▶ Six modified files (don't modify, but **read it**)

- ▶ include/linux/sched.h, kernel/sched.c, kernel/sched_fair.c,
include/linux/syscalls.h,
arch/x86/kernel/syscall_table_32.S,
arch/x86/include/asm/unistd_32.h

- ▶ sched_weighted_rr.c (**incomplete, your job!**)

- enqueue_task_weighted_rr(), dequeue_task_weighted_rr()

- yield_task_weighted_rr(), pick_next_task_weighted_rr()

- task_tick_weighted_rr()

Testing Program

In linux-2.6.32.60\test_weighted_rr\test_weighted_rr.c

- ▶ The test program will first allocate a write buffer with size **b**.
- ▶ Then, the test program will create **n** user threads, each of which will write a unique character (e.g., a) into the buffer over and over.
 - ▶ Note that, every threads will write the same number of characters in to the buffer, based on the buffer size.
- ▶ Moreover, you can assign the scheduling policy, and the weighted_rr_time_slice **t**.

Note that, when dumping the write buffer, the test program will aggregate the consecutive characters into one symbol.

Possible Results

```
▶ ./test_weighted_rr weighted_rr t 5 5000000
```

```
> ./test_weighted_rr weighted_rr 10 5 500000000  
sched_policy: 6, quantum: 10, num_threads: 5, buffer_size: 500000000  
abcdeabcdeabcdeabcdabcbcabcbcababababababababa
```

'e' finish 'd' finish 'c' finish

test_weighted_rr.c

```
for (i = 0; i < num_threads; i++)
{
    ...
    syscall (SYS_weighted_rr_setquantum, quantum);
    pthread_create(&threads[i], &attr, run, (void *)targs);
    quantum*=2;
}
```

Scoring of Project 2

- ▶ Part I: Implementation of a program to invoke FIFO scheduler (30%)
- ▶ Part II: Implementation of the below **FIVE** incomplete functions in “sched_weighted_rr.c” (40%)
 - enqueue_task_weighted_rr(), dequeue_task_weighted_rr()
 - yield_task_weighted_rr(), pick_next_task_weighted_rr()
 - task_tick_weighted_rr()
- ▶ Report (30%)
 - ▶ Your implementation details and results
 - ▶ At most 4 pages (6 pages if you have bonus)
- ▶ Bonus (at most 20%)
 - ▶ Any variation of the scheduling policy

Outline

- ▶ Review: Life of Process
- ▶ Scheduling in Linux
- ▶ Implementation
 - ▶ Part I: Invoke FIFO Scheduler
 - ▶ Part II: Weighted Round Robin Scheduler
- ▶ Submission Rules
- ▶ References

Submission Rules

- ▶ Project deadline: 2017/05/10 23:59
- ▶ Upload to FTP Server
 - ▶ IP: 140.112.28.143
 - ▶ Port:10400
 - ▶ Account: os2017
 - ▶ Password: ktw2017os
- ▶ Be packed as one file named “OSPJ2_Team##_v###.zip ”
 - ▶ Report.pdf
 - ▶ Part1(directory)
 - +--- sched_test.c
 - ▶ Part2(directory)
 - +--- sched_weighted_rr.c
 - ▶ Bonus(directory)
 - +--- your code and any modified files
- ▶ **DO NOT COPY THE HOMEWORK**

Contact TAs

- ▶ If you have any problem about the projects, you can contact TAs by the following ways:
- ▶ Facebook: NTU CSIE OS2017 Group
 - ▶ <https://www.facebook.com/groups/380026635712953/>
- ▶ E-mail:
 - ▶ Han-Yi Lin: d03922006@csie.ntu.edu.tw
 - ▶ Chun-Feng Wu: tom.cfwu@gmail.com
 - ▶ Yu-Chen Lin: f04922077@csie.ntu.edu.tw
 - ▶ Chih-Hsuan Yen: r04922036@csie.ntu.edu.tw

References

- ▶ Reference Book

- ▶ Professional Linux® Kernel Architecture, Wolfgang
Mauerer, Wiley Publishing, Inc.

- ▶ Process Scheduling

- ▶ <https://www.cs.rutgers.edu/~pxk/416/notes/07-scheduling.html>