

Tracing codes TIPS (I)

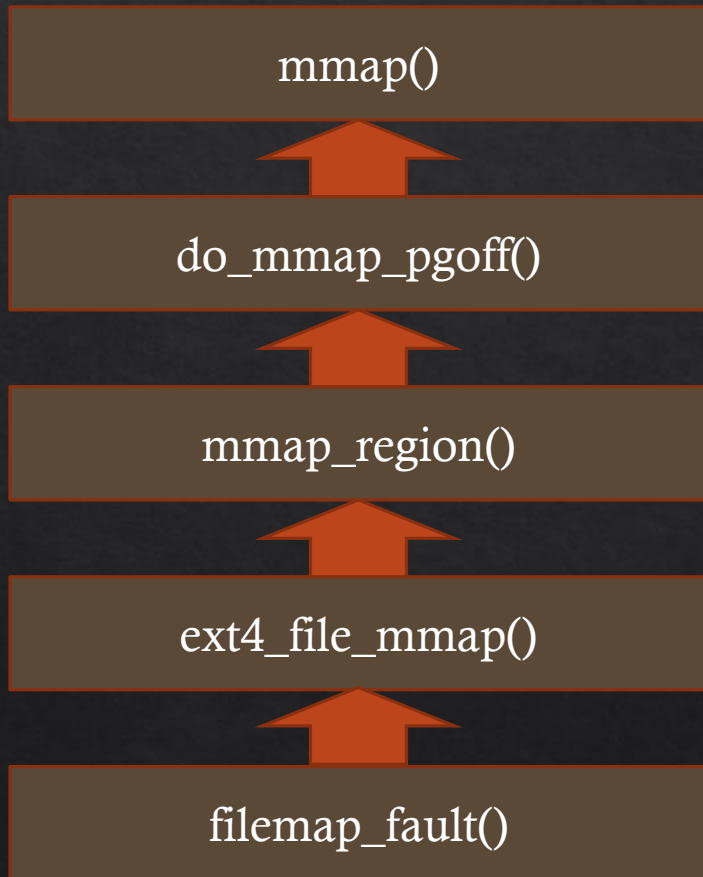
◇ Steps

- ◇ Google for a function name
- ◇ Bottom-up backtracking
- ◇ Top-down analyzing

◇ Tools

- ◇ `grep -nr`: generic text search
- ◇ ctags: semantic-aware search
- ◇ LXR – Linux Cross Reference
<http://lxr.linux.no/#linux+v2.6.32.60/>

Readahead flow



Tracing codes TIPS (II)


❖ Object Oriented Programming in C

```
/* fs/ext4/file.c */  
static const struct vm_operations_struct ext4_file_vm_ops = {  
    .fault          = filemap_fault,  
    .page_mkwrite   = ext4_page_mkwrite,  
};
```

```
/* mm/filemap.c */  
const struct vm_operations_struct generic_file_vm_ops = {  
    .fault          = filemap_fault,  
};
```

`fault` is a “virtual function”

Readahead algorithm



No readahead
(Pure demand paging)

Better for SSD:
reduce unnecessary
works

Full readahead

Better for HDD:
amortize seek
time cost

Readahead algorithm

◇ How many pages are read?

```
/* mm/filemap.c */
```

```
page_cache_async_readahead(mapping, ra, file, page,  
                             offset, ra->ra_pages);
```

```
/* mm/readahead.c */
```

```
ra->ra_pages = mapping->backing_dev_info->ra_pages;
```

```
/* mm/backing-dev.c */
```

```
struct backing_dev_info default_backing_dev_info = {  
    .ra_pages = VM_MAX_READAHEAD * 1024 / PAGE_CACHE_SIZE,  
};
```


Project Goals

◆ Code reading

- ◆ Part I: How `filemap_fault()` is set as the page fault handler when `mmap()` is called
- ◆ Part II: How and when the readahead algorithm takes place when `filemap_fault()` is invoked

◆ Revise readahead algorithm

- ◆ Any change that reduces the time between “page fault test program starts !” and “page fault test program ends !”
- ◆ Percentage of reduction varies from machine to machine. Any number larger than 0 is fine. As a reference: 6% on SSD and > 10% on HDD